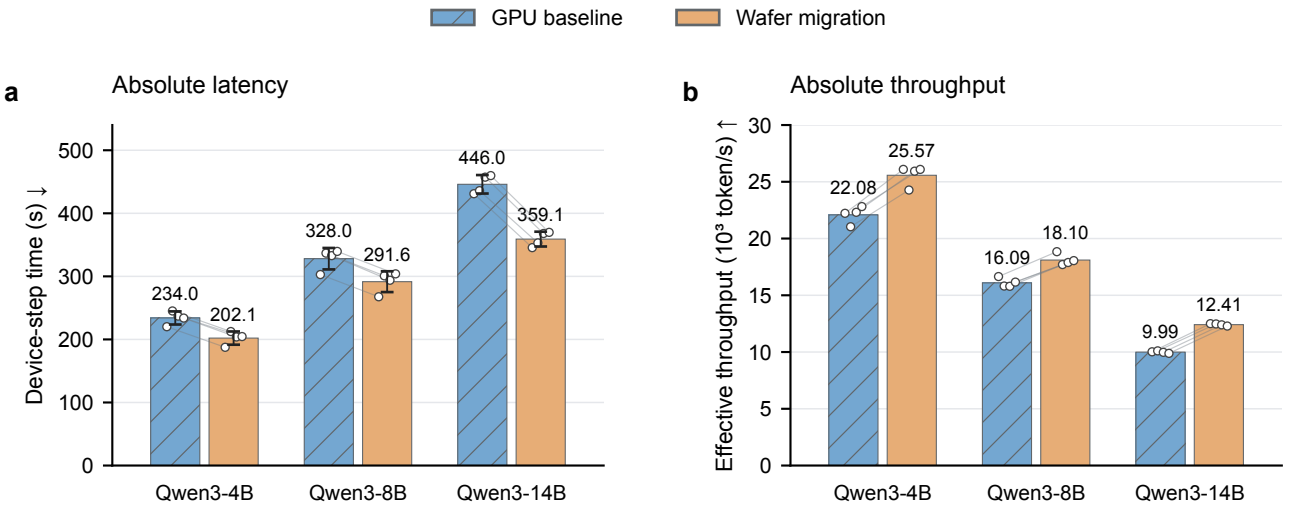

Time bars: mean ± SD. Throughput bars: total tokens / total time. Points: 4 matched windows.

Effective tokens = prompt + response, counted once per trajectory. 192 prompts × 8 answers per step.

Current uncalibrated simulation; window variation is not a hardware-model error bound.